

Module_1_intro_dplyr

Elyse, Jeanne & Sheila

Table of contents

Welcome & Setup	2
Learning Objectives (Module 1)	2
The Tools We'll Use	2
Load Your Tools	3
Step 1: Create a Directory	5
Step 2: Download the Dataset	5
What Makes Big Data Big?	6
Core Concepts: The 3 Vs	6
The Three V's Framework	6
B. dplyr Refresher	7
Why this refresher?	7
The dplyr Mindset: Think in Verbs	7
Meet Your Practice Dataset: starwars	8
Use Case Dataset: starwars	8
filter(): Keep only rows that meet a condition	9
select(): Pick specific columns	13
arrange(): Sort rows	17
mutate(): Create or Modify names	19
group_by()	22
summarise()	23
Why this matters for arrow	25
Typical dplyr Workflow (Local or Scalable)	25
Step-by-Step Pattern:	25
Common Mistakes and Solutions	26
1. Forgetting the Pipe	26
2. Not Handling Missing Values	27
3. Forgetting to Ungroup	27
4. Incorrect Logical Operators	27

Try It Yourself: Challenges	28
Beginner Challenge (5-10 minutes)	28
Intermediate Challenge (10-15 minutes)	30
Why These Challenges Matter	34
What's Coming Next	35

Welcome & Setup

Welcome to our short course! Today we'll tackle a common pain point in R: **what happens when your data is just too big?**

But first, let's get everyone set up and comfortable.

Learning Objectives (Module 1)

By the end of this module, you will be able to:

- Master the 6 core dplyr verbs: `filter()`, `select()`, `arrange()`, `mutate()`, `group_by()`, and `summarise()`
- Build multi-step data transformation pipelines using the pipe operator (`|>`)
- Apply the “filter early, select early” optimization strategy
- Construct grouped summaries and aggregations
- Understand the analytical thinking process that scales to big data
- Practice troubleshooting common dplyr errors

The Tools We'll Use

The tools we're introducing today **each play a specific role** in helping us work with larger-than-memory datasets — and we can combine them to build scalable workflows.

Tool	Purpose
arrow	(Columnar storage + lazy reading) → Efficiently reads large datasets without loading the whole file into memory . Supports fast filtering and streaming from disk.
DBI	(Database connection interface) → Provides a common language to connect R to databases . DuckDB uses it to talk to R.

Tool	Purpose
duckdb	(Fast in-process SQL database) → Allows you to query large datasets using SQL syntax inside R . Works especially well for joins, aggregations, and window functions.
dplyr	(User-friendly data wrangling) → Offers an intuitive, readable grammar for filtering, summarizing, and transforming data . It's our main pipeline tool.
dbplyr	(Bridge between dplyr and databases) → Translates dplyr code into SQL automatically . Lets you use dplyr pipelines on database tables (like DuckDB) without writing raw SQL.

•

Load Your Tools

Let's load the packages we'll need today.

```
# Install and load required packages
required_packages <- c("tidyverse", "arrow", "duckdb", "DBI", "dbplyr")

# Install missing packages
for (pkg in required_packages) {
  if (!requireNamespace(pkg, quietly = TRUE)) {
    install.packages(pkg)
  }
}

# Load all packages
for (pkg in required_packages) {
  library(pkg, character.only = TRUE)
}
```

Warning: package 'tidyverse' was built under R version 4.3.3

Warning: package 'tibble' was built under R version 4.3.3

Warning: package 'tidyr' was built under R version 4.3.3

Warning: package 'readr' was built under R version 4.3.3

Warning: package 'purrr' was built under R version 4.3.3

Warning: package 'dplyr' was built under R version 4.3.3

Warning: package 'stringr' was built under R version 4.3.3

Warning: package 'forcats' was built under R version 4.3.3

Warning: package 'lubridate' was built under R version 4.3.3

-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --

v dplyr 1.1.4 v readr 2.1.5

v forcats 1.0.0 v stringr 1.5.1

v ggplot2 3.5.2 v tibble 3.2.1

v lubridate 1.9.3 v tidyr 1.3.1

v purrr 1.0.2

-- Conflicts ----- tidyverse_conflicts() --

x dplyr::filter() masks stats::filter()

x dplyr::lag() masks stats::lag()

i Use the conflicted package (<<http://conflicted.r-lib.org/>>) to force all conflicts to become

Warning: package 'arrow' was built under R version 4.3.3

Attaching package: 'arrow'

The following object is masked from 'package:lubridate':

duration

The following object is masked from 'package:utils':

timestamp

Warning: package 'duckdb' was built under R version 4.3.3

Loading required package: DBI

Warning: package 'DBI' was built under R version 4.3.3

Warning: package 'dbplyr' was built under R version 4.3.3

Attaching package: 'dbplyr'

The following objects are masked from 'package:dplyr':

ident, sql

For this shortcourse we will be using the [Seattle Library Checkouts Dataset](#) if you didn't do it in the pre-work then please do it now. We will talk about this in the Module 2- Arrow.

Step 1: Create a Directory

First, let's create a special folder to store our data:

```
# Create a "data" directory if it doesn't exist already
# Using showWarnings = FALSE to suppress warning if directory already exists

dir.create("data", showWarnings = FALSE)
```

Step 2: Download the Dataset

Now for the fun part! We'll download the Seattle Library dataset (9GB).

Important: This is a 9GB file, so:

- Make sure you have enough disk space

```
# Download Seattle library checkout dataset:

# 1. Fetch data from AWS S3 bucket URL
# 2. Save to local data directory
# 3. Use resume = TRUE to allow continuing interrupted downloads

curl::multi_download("https://r4ds.s3.us-west-2.amazonaws.com/seattle-library-checkouts.csv")
```

What Makes Big Data Big?

Tip: You might see errors like “cannot allocate vector of size...”

Example:

```
# simulation
big_data_test <- tibble(x = rnorm(50000000))

big_data_test
```

```
# A tibble: 50,000,000 x 1
      x
  <dbl>
1  0.237
2 -0.195
3  1.57
4 -0.192
5  0.521
6 -0.307
7 -0.478
8 -0.751
9 -0.280
10 0.877
# i 49,999,990 more rows
```

```
# Attempt to simulate a "too large" allocation
big_data_test <- tibble(x = rnorm(1e9)) # 1 billion rows
```

Core Concepts: The 3 Vs

The Three V's Framework

Volume: The size of your data exceeds your computer's RAM

- Example: A 5GB CSV file on a laptop with 8GB RAM

Velocity: Data is generated faster than you can process it

- Example: Real-time sensor data, streaming transactions

Variety: Different data types and structures

- Example: Combining CSV files, JSON logs, and database tables

Interactive discussion:

- What's the largest dataset you've tried to analyze in R?
- When has R crashed on you? What error messages did you see?
- How do you currently handle data that's "too big"?

Common strategies people use (and their limitations):

- **Sampling:** Loses information and representativeness
- **Chunking:** Complex to manage and prone to errors
- **Giving up:** Not a solution!
- **Today's approach:** Use the right tools for the job

B. dplyr Refresher

Why this refresher?

Before we dive into **arrow** and **bigger-than-memory data**, let's quickly revisit **dplyr**, the backbone of tidy data manipulation in R.

Experience check:

- Daily dplyr users (you'll be our helpers!)
- Occasional users (perfect timing for a refresher!)
- New to dplyr (you're in for a treat!)

The dplyr Mindset: Think in Verbs

Traditional R vs. dplyr Approach

Traditional R thinking:

```
# Multiple objects, hard to follow
subset_data <- subset(starwars, species == "Human")
selected_data <- subset_data[, c("name", "height", "mass")]
ordered_data <- selected_data[order(selected_data$height, decreasing = TRUE), ]
```

dplyr thinking:

```
# One pipeline, easy to read
starwars |>
  filter(species == "Human") |>
  select(name, height, mass) |>
  arrange(desc(height))
```

```
# A tibble: 35 x 3
   name          height mass
   <chr>         <int> <dbl>
1 Darth Vader      202  136
2 Qui-Gon Jinn     193   89
3 Dooku            193   80
4 Bail Prestor Organa 191  NA
5 Anakin Skywalker  188   84
6 Mace Windu       188   84
7 Raymus Antilles   188   79
8 Padmé Amidala     185   45
9 Biggs Darklighter 183   84
10 Boba Fett        183  78.2
# i 25 more rows
```

Before we scale up to big data, let's review the `dplyr` verbs that will be our building blocks.

Rule of thumb:

- Small datasets (< 1MB): Either approach works
- Medium datasets (1MB - 100MB): `dplyr` is cleaner
- Large datasets (100MB+): `dplyr` pipeline thinking is essential

Meet Your Practice Dataset: `starwars`

Built in dataset perfect for learning

Use Case Dataset: `starwars`

We'll work with the built-in `starwars` tibble.

```
#check out the first 5 rows
starwars |>
  head()
```



```
# A tibble: 6 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sky~   172    77 blond      fair        blue         19  male  mascu~
2 C-3PO      167    75 <NA>      gold        yellow       112  none  mascu~
3 R2-D2      96     32 <NA>      white, bl~  red          33  none  mascu~
4 Darth Va~  202   136 none      white       yellow       41.9  male  mascu~
5 Leia Org~  150    49 brown     light       brown        19  fema~  femin~
6 Owen Lars  178   120 brown, gr~ light       blue         52  male  mascu~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

filter(): Keep only rows that meet a condition

```
# Single condition
starwars |>
  filter(species == "Human")
```

```
# A tibble: 35 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~   172    77 blond      fair        blue         19  male  mascu~
2 Darth V~   202   136 none      white       yellow       41.9  male  mascu~
3 Leia Or~   150    49 brown     light       brown        19  fema~  femin~
4 Owen La~   178   120 brown, gr~ light       blue         52  male  mascu~
5 Beru Wh~   165    75 brown     light       blue         47  fema~  femin~
6 Biggs D~   183    84 black     light       brown        24  male  mascu~
7 Obi-Wan~   182    77 auburn, w~ fair        blue-gray    57  male  mascu~
8 Anakin ~   188    84 blond     fair        blue         41.9  male  mascu~
9 Wilhuff~   180    NA auburn, g~ fair        blue         64  male  mascu~
10 Han Solo   180    80 brown     fair        brown        29  male  mascu~
# i 25 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

```
# Multiple conditions (AND)
# Find all humans from Tatooine
starwars |>
  filter(species == "Human", homeworld == "Tatooine")
```

```
# A tibble: 8 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sky~   172    77 blond      fair        blue        19   male masculi~
2 Darth Va~   202   136 none       white       yellow      41.9 male masculi~
3 Owen Lars   178   120 brown, gr~ light       blue        52   male masculi~
4 Beru Whi~   165    75 brown      light       blue        47   fema~ femin~
5 Biggs Da~   183    84 black      light       brown       24   male masculi~
6 Anakin S~   188    84 blond      fair        blue        41.9 male masculi~
7 Shmi Sky~   163    NA black      fair        brown       72   fema~ femin~
8 Cliegg L~   183    NA brown      fair        blue        82   male masculi~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

```
# Multiple conditions (OR)
starwars |>
  filter(species == "Human" | species == "Droid")
```

```
# A tibble: 41 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~   172    77 blond      fair        blue        19   male masculi~
2 C-3PO     167    75 <NA>      gold        yellow     112   none masculi~
3 R2-D2      96    32 <NA>      white, bl~ red        33   none masculi~
4 Darth V~   202   136 none       white       yellow     41.9 male masculi~
5 Leia Or~   150    49 brown      light       brown       19   fema~ femin~
6 Owen La~   178   120 brown, gr~ light       blue        52   male masculi~
7 Beru Wh~   165    75 brown      light       blue        47   fema~ femin~
8 R5-D4      97    32 <NA>      white, red red        NA    none masculi~
9 Biggs D~   183    84 black      light       brown       24   male masculi~
10 Obi-Wan~   182    77 auburn, w~ fair        blue-gray   57   male masculi~
# i 31 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

```
# Numeric conditions
starwars |>
  filter(height > 180, mass < 100)
```

```
# A tibble: 22 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
```

```

      <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Biggs D~      183  84   black      light      brown        24   male  mascu~
2 Obi-Wan~      182  77   auburn, w~ fair      blue-gray    57   male  mascu~
3 Anakin ~      188  84   blond      fair      blue        41.9 male  mascu~
4 Boba Fe~      183  78.2 black      fair      brown        31.5 male  mascu~
5 Qui-Gon~      193  89   brown      fair      blue        92   male  mascu~
6 Nute Gu~      191  90   none      mottled g~ red        NA    male  mascu~
7 Padmé A~      185  45   brown      light      brown        46   fema~ femin~
8 Jar Jar~      196  66   none      orange     orange        52   male  mascu~
9 Roos Ta~      224  82   none      grey       orange        NA    male  mascu~
10 Mace Wi~      188  84   none      dark       brown        72   male  mascu~
# i 12 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

```

# Handle missing values
starwars |>
  filter(!is.na(height))

```

```

# A tibble: 81 x 14
   name      height  mass hair_color skin_color eye_color birth_year sex  gender
   <chr>    <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~    172    77 blond      fair      blue        19   male  mascu~
2 C-3PO      167    75 <NA>      gold      yellow     112   none  mascu~
3 R2-D2       96    32 <NA>      white, bl~ red        33   none  mascu~
4 Darth V~   202   136 none      white      yellow     41.9 male  mascu~
5 Leia Or~   150    49 brown      light      brown        19   fema~ femin~
6 Owen La~   178   120 brown, gr~ light      blue        52   male  mascu~
7 Beru Wh~   165    75 brown      light      blue        47   fema~ femin~
8 R5-D4       97    32 <NA>      white, red red        NA    none  mascu~
9 Biggs D~   183    84 black      light      brown        24   male  mascu~
10 Obi-Wan~   182    77 auburn, w~ fair      blue-gray    57   male  mascu~
# i 71 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

Advanced filter()ing

```

# Using %in% for multiple values
starwars |>
  filter(homeworld %in% c("Tatooine", "Naboo", "Alderaan"))

```

```
# A tibble: 24 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>    <int> <dbl> <chr>      <chr>    <chr>      <dbl> <chr> <chr>
1 Luke Sk~    172    77 blond      fair      blue        19  male  mascu~
2 C-3P0      167    75 <NA>       gold      yellow     112  none  mascu~
3 R2-D2       96    32 <NA>       white, bl~ red        33  none  mascu~
4 Darth V~   202   136 none       white      yellow    41.9  male  mascu~
5 Leia Or~   150    49 brown      light     brown      19  fema~ femin~
6 Owen La~   178   120 brown, gr~ light     blue       52  male  mascu~
7 Beru Wh~   165    75 brown      light     blue       47  fema~ femin~
8 R5-D4       97    32 <NA>       white, red red        NA  none  mascu~
9 Biggs D~   183    84 black      light     brown      24  male  mascu~
10 Anakin ~   188    84 blond      fair      blue     41.9  male  mascu~
# i 14 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

```
# String matching
starwars |>
  filter(str_detect(name, "Skywalker"))
```

```
# A tibble: 3 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>    <int> <dbl> <chr>      <chr>    <chr>      <dbl> <chr> <chr>
1 Luke Sky~    172    77 blond      fair      blue        19  male  mascu~
2 Anakin S~   188    84 blond      fair      blue     41.9  male  mascu~
3 Shmi Sky~   163    NA black      fair      brown       72  fema~ femin~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

```
# Complex logical conditions
starwars |>
  filter((species == "Human" & height > 175) | species == "Droid")
```

```
# A tibble: 25 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>    <int> <dbl> <chr>      <chr>    <chr>      <dbl> <chr> <chr>
1 C-3P0      167    75 <NA>       gold      yellow     112  none  mascu~
2 R2-D2       96    32 <NA>       white, bl~ red        33  none  mascu~
3 Darth V~   202   136 none       white      yellow    41.9  male  mascu~
4 Owen La~   178   120 brown, gr~ light     blue       52  male  mascu~
```

```

5 R5-D4          97      32 <NA>      white, red red          NA    none  mascu~
6 Biggs D~      183      84 black      light      brown          24    male  mascu~
7 Obi-Wan~      182      77 auburn, w~ fair      blue-gray       57    male  mascu~
8 Anakin ~      188      84 blond      fair      blue           41.9  male  mascu~
9 Wilhuff~      180      NA auburn, g~ fair      blue           64    male  mascu~
10 Han Solo      180      80 brown      fair      brown          29    male  mascu~
# i 15 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

select(): Pick specific columns

```

# Select specific columns
starwars |>
  select(name, height, mass)

```

```

# A tibble: 87 x 3
   name          height  mass
  <chr>         <int> <dbl>
1 Luke Skywalker    172    77
2 C-3P0             167    75
3 R2-D2              96    32
4 Darth Vader      202   136
5 Leia Organa      150    49
6 Owen Lars        178   120
7 Beru Whitesun Lars 165    75
8 R5-D4              97    32
9 Biggs Darklighter 183    84
10 Obi-Wan Kenobi   182    77
# i 77 more rows

```

```

# Select ranges
starwars |>
  select(name:mass)

```

```

# A tibble: 87 x 3
   name          height  mass
  <chr>         <int> <dbl>
1 Luke Skywalker    172    77
2 C-3P0             167    75

```

```

3 R2-D2          96    32
4 Darth Vader   202   136
5 Leia Organa   150    49
6 Owen Lars     178   120
7 Beru Whitesun 165    75
8 R5-D4         97    32
9 Biggs Darklighter 183   84
10 Obi-Wan Kenobi 182   77
# i 77 more rows

```

```

# Exclude columns
starwars |>
  select(-films, -vehicles, -starships)

```

```

# A tibble: 87 x 11
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~    172    77 blond      fair        blue        19    male masculi~
2 C-3P0      167    75 <NA>      gold        yellow     112    none masculi~
3 R2-D2       96    32 <NA>      white, bl~ red         33    none masculi~
4 Darth V~   202   136 none      white      yellow     41.9  male masculi~
5 Leia Or~   150    49 brown     light      brown       19    fema~ femini~
6 Owen La~   178   120 brown, gr~ light      blue       52    male masculi~
7 Beru Wh~   165    75 brown     light      blue       47    fema~ femini~
8 R5-D4       97    32 <NA>      white, red red         NA    none masculi~
9 Biggs D~   183    84 black     light      brown       24    male masculi~
10 Obi-Wan~   182    77 auburn, w~ fair      blue-gray   57    male masculi~
# i 77 more rows
# i 2 more variables: homeworld <chr>, species <chr>

```

```

# Select by pattern
starwars |>
  select(starts_with("s")) # species, skin_color, starships

```

```

# A tibble: 87 x 4
  skin_color sex  species starships
  <chr>      <chr> <chr>    <list>
1 fair      male  Human   <chr [2]>
2 gold      none  Droid   <chr [0]>
3 white, blue none  Droid   <chr [0]>
4 white     male  Human   <chr [1]>

```

```

5 light      female Human <chr [0]>
6 light      male   Human <chr [0]>
7 light      female Human <chr [0]>
8 white, red none   Droid <chr [0]>
9 light      male   Human <chr [1]>
10 fair      male   Human <chr [5]>
# i 77 more rows

```

```

starwars |>
  select(ends_with("color")) # hair_color, skin_color, eye_color

```

```

# A tibble: 87 x 3
  hair_color skin_color eye_color
  <chr>      <chr>      <chr>
1 blond     fair        blue
2 <NA>      gold        yellow
3 <NA>      white, blue red
4 none      white       yellow
5 brown     light       brown
6 brown, grey light      blue
7 brown     light      blue
8 <NA>      white, red  red
9 black     light      brown
10 auburn, white fair      blue-gray
# i 77 more rows

```

```

starwars |>
  select(contains("_")) # hair_color, skin_color, eye_color, birth_year

```

```

# A tibble: 87 x 4
  hair_color skin_color eye_color birth_year
  <chr>      <chr>      <chr>      <dbl>
1 blond     fair        blue         19
2 <NA>      gold        yellow      112
3 <NA>      white, blue red         33
4 none      white       yellow     41.9
5 brown     light       brown        19
6 brown, grey light      blue        52
7 brown     light      blue        47
8 <NA>      white, red  red         NA
9 black     light      brown        24

```

```
10 auburn, white fair          blue-gray          57
# i 77 more rows
```

Advanced Select() ing

```
# Rename while selecting
starwars |>
  select(character_name = name, height_cm = height)
```

```
# A tibble: 87 x 2
  character_name height_cm
  <chr>          <int>
1 Luke Skywalker    172
2 C-3PO             167
3 R2-D2              96
4 Darth Vader       202
5 Leia Organa       150
6 Owen Lars         178
7 Beru Whitesun Lars 165
8 R5-D4              97
9 Biggs Darklighter 183
10 Obi-Wan Kenobi    182
# i 77 more rows
```

```
# Select and reorder
starwars |>
  select(name, species, everything()) # name and species first, then everything else
```

```
# A tibble: 87 x 14
  name      species height mass hair_color skin_color eye_color birth_year sex
  <chr>    <chr>    <int> <dbl> <chr>    <chr>    <chr>    <dbl> <chr>
1 Luke S~ Human      172    77 blond    fair     blue      19    male
2 C-3PO    Droid      167    75 <NA>    gold     yellow    112    none
3 R2-D2    Droid       96    32 <NA>    white, bl~ red       33    none
4 Darth ~ Human     202   136 none     white    yellow    41.9  male
5 Leia O~ Human     150    49 brown    light    brown     19    fema~
6 Owen L~ Human     178   120 brown, gr~ light    blue      52    male
7 Beru W~ Human     165    75 brown    light    blue      47    fema~
8 R5-D4    Droid       97    32 <NA>    white, red red       NA    none
9 Biggs ~ Human     183    84 black    light    brown     24    male
10 Obi-Wa~ Human     182    77 auburn, w~ fair     blue-gray 57    male
```



```
# i 77 more rows
# i 5 more variables: gender <chr>, homeworld <chr>, films <list>,
#   vehicles <list>, starships <list>
```

arrange(): Sort rows

```
# Sort ascending (default)
starwars |>
  arrange(height)
```

```
# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Yoda        66    17 white      green      brown          896 male  mascu~
2 Ratts T~    79    15 none      grey, blue unknown          NA male  mascu~
3 Wicket ~    88    20 brown     brown      brown           8 male  mascu~
4 Dud Bolt    94    45 none      blue, grey yellow          NA male  mascu~
5 R2-D2       96    32 <NA>      white, bl~ red           33 none  mascu~
6 R4-P17      96    NA none      silver, r~ red, blue          NA none  femin~
7 R5-D4       97    32 <NA>      white, red red           NA none  mascu~
8 Sebulba    112    40 none      grey, red  orange          NA male  mascu~
9 Gasgano     122    NA none      white, bl~ black          NA male  mascu~
10 Watto     137    NA black     blue, grey yellow          NA male  mascu~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

```
# Sort descending
starwars |>
  arrange(desc(height))
```

```
# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Yarael ~    264    NA none      white      yellow          NA male  mascu~
2 Tarfful    234   136 brown     brown      blue           NA male  mascu~
3 Lama Su    229    88 none      grey      black          NA male  mascu~
4 Chewbac~   228   112 brown     unknown    blue          200 male  mascu~
5 Roos Ta~   224    82 none      grey      orange          NA male  mascu~
6 Grievous   216   159 none      brown, wh~ green, y~          NA male  mascu~
```

```

7 Taun We      213    NA none      grey      black      NA   fema~ femin~
8 Rugor N~     206    NA none      green      orange     NA   male  mascu~
9 Tion Me~     206    80 none      grey      black      NA   male  mascu~
10 Darth V~    202   136 none      white      yellow     41.9 male  mascu~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

```

# Multiple sort columns
starwars |>
  arrange(species, desc(height))

```

```

# A tibble: 87 x 14
   name      height  mass hair_color skin_color eye_color birth_year sex  gender
   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Ratts T~      79    15 none      grey, blue unknown      NA male  mascu~
2 Dexter ~     198   102 none      brown      yellow      NA male  mascu~
3 Ki-Adi~~     198    82 white      pale      yellow      92 male  mascu~
4 Mas Ame~     196    NA none      blue      blue      NA male  mascu~
5 Zam Wes~     168    55 blonde    fair, gre~ yellow      NA fema~ femin~
6 IG-88        200   140 none      metal      red        15 none  mascu~
7 C-3PO        167    75 <NA>      gold      yellow     112 none  mascu~
8 R5-D4         97    32 <NA>      white, red red        NA none  mascu~
9 R2-D2         96    32 <NA>      white, bl~ red        33 none  mascu~
10 R4-P17        96    NA none      silver, r~ red, blue   NA none  femin~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

```

# Handle missing values
starwars |>
  arrange(desc(height), na.last = TRUE)

```

```

# A tibble: 87 x 14
   name      height  mass hair_color skin_color eye_color birth_year sex  gender
   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Yarael ~     264    NA none      white      yellow      NA   male  mascu~
2 Tarfful      234   136 brown      brown      blue      NA   male  mascu~
3 Lama Su      229    88 none      grey      black      NA   male  mascu~
4 Chewbac~     228   112 brown      unknown    blue     200   male  mascu~
5 Roos Ta~     224    82 none      grey      orange     NA   male  mascu~

```

```

6 Grievous      216    159 none      brown, wh~ green, y~      NA    male  mascu~
7 Taun We       213      NA none      grey       black          NA    fema~ femin~
8 Rugor N~      206      NA none      green      orange         NA    male  mascu~
9 Tion Me~      206      80 none      grey       black          NA    male  mascu~
10 Darth V~     202     136 none      white      yellow         41.9 male  mascu~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

mutate(): Create or Modify names

```

# Create new column
starwars |>
  mutate(height_m = height / 100)

```

```

# A tibble: 87 x 15
   name      height  mass hair_color skin_color eye_color birth_year sex  gender
   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~    172     77 blond      fair       blue        19    male  mascu~
2 C-3PO      167     75 <NA>      gold       yellow      112    none  mascu~
3 R2-D2       96     32 <NA>      white, bl~ red         33    none  mascu~
4 Darth V~   202    136 none      white      yellow      41.9   male  mascu~
5 Leia Or~   150     49 brown      light      brown       19    fema~ femin~
6 Owen La~   178    120 brown, gr~ light      blue        52    male  mascu~
7 Beru Wh~   165     75 brown      light      blue        47    fema~ femin~
8 R5-D4       97     32 <NA>      white, red red         NA     none  mascu~
9 Biggs D~   183     84 black      light      brown       24    male  mascu~
10 Obi-Wan~  182     77 auburn, w~ fair       blue-gray    57    male  mascu~
# i 77 more rows
# i 6 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>, height_m <dbl>

```

```

# Multiple new columns
starwars |>
  mutate(
    height_m = height / 100,
    bmi = mass / (height_m^2)
  )

```

```

# A tibble: 87 x 16

```

```

  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~   172    77 blond      fair        blue        19    male masculin~
2 C-3PO      167    75 <NA>      gold        yellow      112   none masculin~
3 R2-D2       96    32 <NA>      white, bl~  red        33    none masculin~
4 Darth V~   202   136 none      white       yellow      41.9  male masculin~
5 Leia Or~   150    49 brown     light       brown       19    fema~ feminin~
6 Owen La~   178   120 brown, gr~ light       blue        52    male masculin~
7 Beru Wh~   165    75 brown     light       blue        47    fema~ feminin~
8 R5-D4       97    32 <NA>      white, red  red        NA     none masculin~
9 Biggs D~   183    84 black     light       brown       24    male masculin~
10 Obi-Wan~  182    77 auburn, w~ fair        blue-gray   57    male masculin~
# i 77 more rows
# i 7 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>, height_m <dbl>, bmi <dbl>

```

```

# Modify existing column
starwars |>
  mutate(name = str_to_upper(name))

```

```

# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 LUKE SK~   172    77 blond      fair        blue        19    male masculin~
2 C-3PO      167    75 <NA>      gold        yellow      112   none masculin~
3 R2-D2       96    32 <NA>      white, bl~  red        33    none masculin~
4 DARTH V~   202   136 none      white       yellow      41.9  male masculin~
5 LEIA OR~   150    49 brown     light       brown       19    fema~ feminin~
6 OWEN LA~   178   120 brown, gr~ light       blue        52    male masculin~
7 BERU WH~   165    75 brown     light       blue        47    fema~ feminin~
8 R5-D4       97    32 <NA>      white, red  red        NA     none masculin~
9 BIGGS D~   183    84 black     light       brown       24    male masculin~
10 OBI-WAN~  182    77 auburn, w~ fair        blue-gray   57    male masculin~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

Advanced mutate() ions

```

# Conditional mutations
starwars |>

```

```
mutate(
  size_category = case_when(
    height < 100 ~ "Very Short",
    height < 150 ~ "Short",
    height < 180 ~ "Average",
    height >= 180 ~ "Tall",
    TRUE ~ "Unknown"
  )
)
```

```
# A tibble: 87 x 15
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex	gender
	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>	<chr>
1	Luke Sk~	172	77	blond	fair	blue	19	male	mascu~
2	C-3P0	167	75	<NA>	gold	yellow	112	none	mascu~
3	R2-D2	96	32	<NA>	white, bl~	red	33	none	mascu~
4	Darth V~	202	136	none	white	yellow	41.9	male	mascu~
5	Leia Or~	150	49	brown	light	brown	19	fema~	femin~
6	Owen La~	178	120	brown, gr~	light	blue	52	male	mascu~
7	Beru Wh~	165	75	brown	light	blue	47	fema~	femin~
8	R5-D4	97	32	<NA>	white, red	red	NA	none	mascu~
9	Biggs D~	183	84	black	light	brown	24	male	mascu~
10	Obi-Wan~	182	77	auburn, w~	fair	blue-gray	57	male	mascu~

```
# i 77 more rows
```

```
# i 6 more variables: homeworld <chr>, species <chr>, films <list>,
```

```
# vehicles <list>, starships <list>, size_category <chr>
```

```
# Using ifelse for simple conditions
```

```
starwars |>
```

```
mutate(is_tall = ifelse(height > 180, "Tall", "Not Tall"))
```

```
# A tibble: 87 x 15
```

	name	height	mass	hair_color	skin_color	eye_color	birth_year	sex	gender
	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>	<chr>
1	Luke Sk~	172	77	blond	fair	blue	19	male	mascu~
2	C-3P0	167	75	<NA>	gold	yellow	112	none	mascu~
3	R2-D2	96	32	<NA>	white, bl~	red	33	none	mascu~
4	Darth V~	202	136	none	white	yellow	41.9	male	mascu~
5	Leia Or~	150	49	brown	light	brown	19	fema~	femin~
6	Owen La~	178	120	brown, gr~	light	blue	52	male	mascu~
7	Beru Wh~	165	75	brown	light	blue	47	fema~	femin~

```

8 R5-D4          97      32 <NA>      white, red red          NA    none  mascu~
9 Biggs D~      183      84 black      light      brown          24    male  mascu~
10 Obi-Wan~     182      77 auburn, w~ fair        blue-gray       57    male  mascu~
# i 77 more rows
# i 6 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>, is_tall <chr>

```

group_by()

```

# Group by single variable
starwars |>
  group_by(species)

```

```

# A tibble: 87 x 14
# Groups:   species [38]
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~    172    77 blond      fair        blue        19    male  mascu~
2 C-3PO      167    75 <NA>      gold        yellow      112    none  mascu~
3 R2-D2       96    32 <NA>      white, bl~ red         33    none  mascu~
4 Darth V~   202   136 none       white       yellow      41.9  male  mascu~
5 Leia Or~   150    49 brown      light       brown       19    fema~ femin~
6 Owen La~   178   120 brown, gr~ light       blue        52    male  mascu~
7 Beru Wh~   165    75 brown      light       blue        47    fema~ femin~
8 R5-D4       97    32 <NA>      white, red red          NA    none  mascu~
9 Biggs D~   183    84 black      light       brown       24    male  mascu~
10 Obi-Wan~  182    77 auburn, w~ fair        blue-gray    57    male  mascu~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

```

# Group by multiple variables
starwars |>
  group_by(species, homeworld)

```

```

# A tibble: 87 x 14
# Groups:   species, homeworld [57]
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Luke Sk~    172    77 blond      fair        blue        19    male  mascu~

```

```

2 C-3P0      167    75 <NA>      gold      yellow      112    none  mascu~
3 R2-D2      96     32 <NA>      white, bl~ red       33     none  mascu~
4 Darth V~   202   136 none      white      yellow     41.9   male  mascu~
5 Leia Or~   150    49 brown     light      brown      19     fema~ femin~
6 Owen La~   178   120 brown, gr~ light      blue       52     male  mascu~
7 Beru Wh~   165    75 brown     light      blue       47     fema~ femin~
8 R5-D4      97     32 <NA>      white, red red       NA     none  mascu~
9 Biggs D~   183    84 black     light      brown      24     male  mascu~
10 Obi-Wan~   182    77 auburn, w~ fair       blue-gray   57     male  mascu~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

summarise()

```

# Single summary
starwars |>
  group_by(species) |>
  summarise(avg_height = mean(height, na.rm = TRUE))

```

```

# A tibble: 38 x 2
  species    avg_height
  <chr>         <dbl>
1 Aleena         79
2 Besalisk      198
3 Cerean        198
4 Chagrian      196
5 Clawdite      168
6 Droid        131.
7 Dug           112
8 Ewok           88
9 Geonosian     183
10 Gungan       209.
# i 28 more rows

```

```

# Multiple summaries
starwars |>
  group_by(species) |>
  summarise(
    count = n(),

```

```

    avg_height = mean(height, na.rm = TRUE),
    max_mass = max(mass, na.rm = TRUE),
    .groups = "drop" # Ungroup after summarising
  )

```

Warning: There were 6 warnings in `summarise()`.

The first warning was:

i In argument: `max_mass = max(mass, na.rm = TRUE)`.

i In group 4: `species = "Chagrian"`.

Caused by warning in `max()`:

! no non-missing arguments to max; returning -Inf

i Run `dplyr::last_dplyr_warnings()` to see the 5 remaining warnings.

A tibble: 38 x 4

	species <chr>	count <int>	avg_height <dbl>	max_mass <dbl>
1	Aleena	1	79	15
2	Besalisk	1	198	102
3	Cerean	1	198	82
4	Chagrian	1	196	-Inf
5	Clawdite	1	168	55
6	Droid	6	131.	140
7	Dug	1	112	40
8	Ewok	1	88	20
9	Geonosian	1	183	80
10	Gungan	3	209.	82

i 28 more rows

Advanced Summarize

```

# Conditional summaries
starwars |>
  group_by(species) |>
  summarise(
    count = n(),
    humans_count = sum(species == "Human", na.rm = TRUE),
    avg_height = mean(height, na.rm = TRUE),
    height_range = max(height, na.rm = TRUE) - min(height, na.rm = TRUE),
    .groups = "drop"
  )

```



```
# A tibble: 38 x 5
  species    count humans_count avg_height height_range
  <chr>      <int>      <int>      <dbl>      <int>
1 Aleena         1          0         79          0
2 Besalisk        1          0        198          0
3 Cerean          1          0        198          0
4 Chagrian        1          0        196          0
5 Clawdite        1          0        168          0
6 Droid           6          0       131.         104
7 Dug             1          0        112          0
8 Ewok            1          0         88          0
9 Geonosian       1          0        183          0
10 Gungan         3          0       209.         28
# i 28 more rows
```

Why this matters for arrow

Each of these verbs has an equivalent when working with **arrow-backed data**, allowing you to **scale from local tibbles to massive datasets** without changing your dplyr workflow.

Next up: we'll see how to **read and query large datasets** using **arrow**, keeping the same grammar you already know.

Typical dplyr Workflow (Local or Scalable)

Here's a general-purpose dplyr pipeline workflow that applies to most tidyverse-style data tasks — whether you're working with small data (**tibble**), big local files (**arrow**), or SQL-like queries (**duckdb**):

Step-by-Step Pattern:

1. Read or Connect to the Data

- `read_csv()`, `read_parquet()` for files (we will see this in the next module)
- `open_dataset()` (arrow)
- `dbConnect()` + `tbl()` (duckdb)

2. Initial Filtering

- `filter()` to narrow rows of interest early

3. Select Columns

- `select()` to reduce memory footprint and focus

4. Mutate or Transform

- `mutate()` to derive new columns (e.g., unit conversions, parsing)

5. Group and Summarise

- `group_by()` + `summarise()` for aggregate

6. Arrange or Rank

- `arrange()` or `mutate(rank = ...)` to sort results

7. Join or Bind

- `left_join()`, `bind_rows()` as needed

8. Collect to Memory (next session)

- `collect()` for arrow or duckdb workflows when you're ready to compute

9. Visualize or Write Out

- `ggplot()`, `write_csv()`, or store to `.parquet`, `.csv`, `.duckdb`, etc.

Building Pipelines: The dplyr Way Here's the general pattern we'll use throughout the workshop:

```
data |>                                     # 1. Start with data
  filter(condition) |>                     # 2. Filter early (reduce rows)
  select(relevant_columns) |>             # 3. Select early (reduce columns)
  mutate(new_variables) |>               # 4. Transform as needed
  group_by(grouping_vars) |>             # 5. Group for summaries
  summarise(summary_stats) |>           # 6. Calculate aggregates
  arrange(sorting_column)                # 7. Sort results
```

Common Mistakes and Solutions

1. Forgetting the Pipe

```
# Wrong - breaks the pipeline
starwars |>
  filter(species == "Human")
select(name, height) # This won't work!

# Correct s
tarwars |>
  filter(species == "Human") |>
  select(name, height)
```

2. Not Handling Missing Values

```
# Will return NA
starwars |>
  summarise(avg_height = mean(height))

# Handle NAs explicitly s
tarwars |>
  summarise(avg_height = mean(height, na.rm = TRUE))
```

3. Forgetting to Ungroup

```
# Leaves data grouped (can cause issues later)
grouped_data <- starwars |>
  group_by(species) |>
  summarise(count = n())
# Always ungroup when done
clean_data <- starwars |>
  group_by(species) |>
  summarise(count = n(), .groups = "drop")
```

4. Incorrect Logical Operators

```
# Wrong - this is assignment, not comparison
starwars |>
```

```

filter(species = "Human")

# Correct - use == for comparison
starwars |>
  filter(species == "Human")

# Wrong - can't use && in filter
starwars |>
  filter(height > 180 && mass < 100)

# Correct - use & or separate conditions
starwars |>
  filter(height > 180 & mass < 100) # OR starwars |> filter(height > 180, mass < 100)

```

Try It Yourself: Challenges

Now it's your turn to practice! Work with a partner and help each other or take a brain break!

Beginner Challenge (5-10 minutes)

Goal: Practice the basic dplyr verbs with the `starwars` dataset.

Your mission: Find the tallest character from each homeworld.

Hints:

1. Start with `starwars`
2. Remove rows where `height` is missing
3. Group by `homeworld`
4. Find the maximum height in each group
5. Keep only the name, homeworld, and height columns

```

# Your code here - try before looking at the solution!

# Solution with detailed steps
tallest_by_homeworld <- starwars |>
  # Step 1: Remove rows where height is missing
  filter(!is.na(height)) |>

```

```

# Step 2: Group by homeworld to analyze each separately
group_by(homeworld) |>

# Step 3: For each homeworld, find the character with max height
filter(height == max(height)) |> # Keep only the tallest in each group

# Step 4: Select only the columns we care about
select(name, homeworld, height) |>

# Step 5: Sort by height for easier reading
arrange(desc(height)) |>

# Step 6: Remove grouping for clean output
ungroup()

# Display the results
tallest_by_homeworld

```

```

# A tibble: 49 x 3
  name          homeworld height
  <chr>         <chr>      <int>
1 Yarael Poof   Quermia      264
2 Tarfful       Kashyyyk     234
3 Lama Su       Kamino       229
4 Roos Tarpals  Naboo        224
5 Grievous      Kalee        216
6 Tion Medon    Utapau       206
7 Darth Vader   Tatooine     202
8 IG-88         <NA>         200
9 Ki-Adi-Mundi  Cerea        198
10 Dexter Jettster Ojom        198
# i 39 more rows

```

Stretch goal: Can you also find the shortest character from each homeworld?

```

# Find the shortest character from each homeworld
shortest_by_homeworld <- starwars |>
  filter(!is.na(height)) |>
  group_by(homeworld) |>
  filter(height == min(height)) |> # Change max to min
  select(name, homeworld, height) |>

```

```

arrange(height) |> # Sort ascending instead of descending
ungroup()

shortest_by_homeworld

```

```

# A tibble: 49 x 3
  name          homeworld height
  <chr>         <chr>     <int>
1 Yoda          <NA>         66
2 Ratts Tyerel  Aleen Minor   79
3 Wicket Systri Warrick Endor         88
4 Dud Bolt      Vulpter      94
5 R2-D2         Naboo        96
6 R5-D4         Tatooine      97
7 Sebulba       Malastare    112
8 Gasgano       Troiken     122
9 Watto         Toydaria    137
10 Leia Organa  Alderaan    150
# i 39 more rows

```

Intermediate Challenge (10-15 minutes)

Goal: Build a more complex analytical pipeline.

Your mission: Create a summary report of species diversity across different homeworlds.

Requirements:

1. Count how many different species live on each homeworld
2. Count the total number of characters from each homeworld
3. Calculate the “diversity ratio” (species count / character count)
4. Include the most common species on each homeworld
5. Sort by diversity ratio (most diverse first)

```

# Your code here - try before looking at the solution!

# Step 1: Create the main diversity summary
homeworld_diversity <- starwars |>
  # Remove rows with missing homeworld or species
  filter(!is.na(homeworld), !is.na(species)) |>

```

```

# Group by homeworld to analyze each separately
group_by(homeworld) |>

# Calculate diversity metrics
summarise(
  # Count unique species
  species_count = n_distinct(species),

  # Count total characters
  character_count = n(),

  # Calculate diversity ratio
  diversity_ratio = species_count / character_count,

  # Find most common species (first alphabetically if tied)
  most_common_species = names(sort(table(species), decreasing = TRUE))[1],

  # Keep groups for potential further operations
  .groups = "keep"
) |>

# Remove grouping and sort by diversity ratio
ungroup() |>
arrange(desc(diversity_ratio))

# Display results
homeworld_diversity

```

```

# A tibble: 46 x 5
  homeworld species_count character_count diversity_ratio most_common_species
  <chr>          <int>          <int>          <dbl> <chr>
1 Aleen Minor      1              1              1 Aleena
2 Bespin            1              1              1 Human
3 Cato Neimo~       1              1              1 Neimodian
4 Cerea            1              1              1 Cerean
5 Champala          1              1              1 Chagrian
6 Chandrila         1              1              1 Human
7 Concord Da~       1              1              1 Human
8 Dathomir          1              1              1 Zabrak
9 Dorin             1              1              1 Kel Dor
10 Endor            1              1              1 Ewok

```

```
# i 36 more rows
```

Stretch goals:

- Add average height and mass by homeworld

```
# Enhanced summary with physical characteristics
enhanced_diversity <- starwars |>
  filter(!is.na(homeworld), !is.na(species)) |>
  group_by(homeworld) |>
  summarise(
    species_count = n_distinct(species),
    character_count = n(),
    diversity_ratio = round(species_count / character_count, 3),

    # Physical characteristics
    avg_height = round(mean(height, na.rm = TRUE), 1),
    avg_mass = round(mean(mass, na.rm = TRUE), 1),

    # Most common species
    most_common_species = first(species[which.max(table(species))]),

    .groups = "drop"
  ) |>
  filter(character_count > 1) |>
  arrange(desc(diversity_ratio))

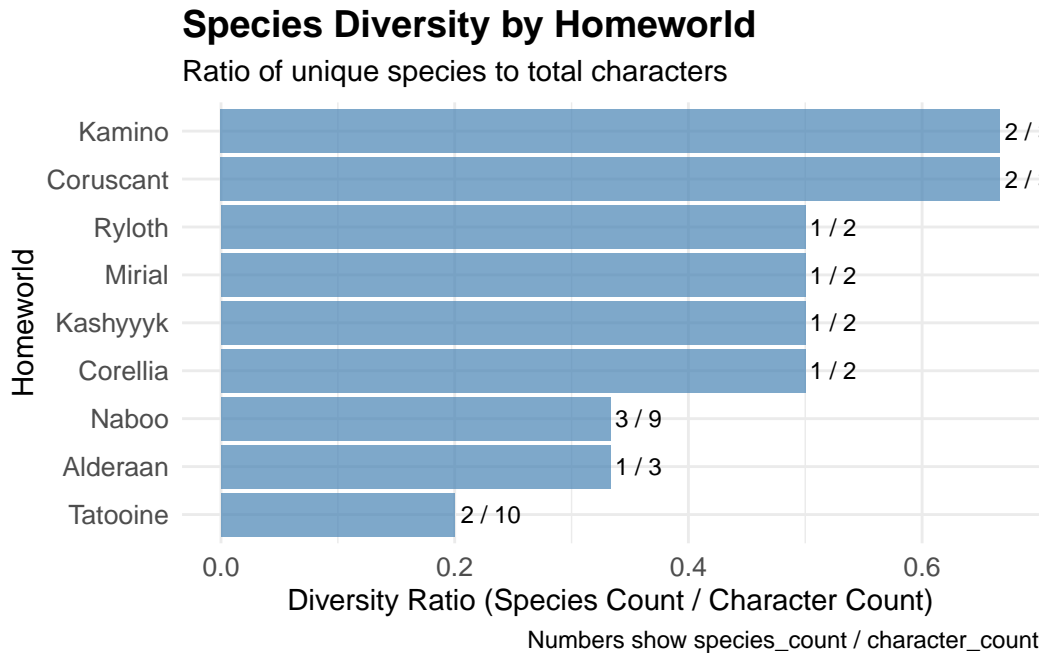
enhanced_diversity
```

```
# A tibble: 9 x 7
  homeworld species_count character_count diversity_ratio avg_height avg_mass
  <chr>          <int>          <int>          <dbl>          <dbl>    <dbl>
1 Coruscant             2              3          0.667          174.     50
2 Kamino                2              3          0.667          208.    83.1
3 Corellia              1              2           0.5           175     78.5
4 Kashyyyk              1              2           0.5           231    124
5 Mirial                1              2           0.5           168     53.1
6 Ryloth                1              2           0.5           179     55
7 Alderaan              1              3          0.333          176.     64
8 Naboo                 3              9          0.333          179.     60
9 Tatooine              2             10           0.2           170.    85.4
# i 1 more variable: most_common_species <chr>
```


- Create a visualization of your results

```
# Visualize diversity across homeworlds
diversity_plot <- homeworld_diversity |>
  filter(character_count > 1) |> # Only homeworlds with multiple characters
  ggplot(aes(x = reorder(homeworld, diversity_ratio), y = diversity_ratio)) +
  geom_col(fill = "steelblue", alpha = 0.7) +
  geom_text(aes(label = paste(species_count, "/", character_count)),
            hjust = -0.1, size = 3) +
  coord_flip() +
  labs(
    title = "Species Diversity by Homeworld",
    subtitle = "Ratio of unique species to total characters",
    x = "Homeworld",
    y = "Diversity Ratio (Species Count / Character Count)",
    caption = "Numbers show species_count / character_count"
  ) +
  theme_minimal() +
  theme(
    plot.title = element_text(size = 14, face = "bold"),
    axis.text = element_text(size = 10)
  )

diversity_plot
```



Why These Challenges Matter

These exercises demonstrate the **exact same thinking process** you'll use with big data:

1. **Filter early** to reduce data size
2. **Select** only what you need
3. **Group and summarize** to aggregate information
4. **Arrange** to present results clearly

The only difference with big data is that we'll add:

- `open_dataset()` instead of using built-in data
- `collect()` at the end to bring results into memory
- `show_query()` to see what's happening behind the scenes

What's Coming Next

In **Module 2**, we'll take these exact same dplyr skills and apply them to:

- A 9GB CSV file (40+ million rows)
- Converting CSV to Parquet for 5x speed improvements
- Processing data that's too large to fit in memory
- Using **arrow** for lazy evaluation and streaming

The promise: Same **dplyr** syntax you just practiced, but on datasets that are 100x larger!

NSF Acknowledgement: This material is based upon work supported by the National Science Foundation under Grant #DGE-2222148. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the National Science Foundation.